opsloom

Menu On this page

Welcome! Opsloom is a tool that allows you to build custom, AI-powered agents that can be used to automate tasks, provide information, or interact with users. These agents are called assistants.

You can give these assistants extra context by supplying them with a repository of documents unique to your organization. These repositories are known as knowledge bases. Using retrieval-augmented generation (RAG) technology, the assistants can be restricted to only answer queries with information from their knowledge bases.

Using tools from Opsloom, you can also empower your assistants to take actions and integrate with your workflows. For example, you could have your assistant query a database and generate a report every so often. Or you could use your assistant to handle incoming customer support calls and route them to the appropriate source.

Read below for some of the concepts central to Opsloom.

Assistants

Assistants are the core building blocks of Opsloom. They are the AI-powered agents that can be used to automate tasks, provide information, or interact with users. You can build your own custom assistants or use the pre-built ones provided by Opsloom.

Large Language Models

A large language model (LLM) is a type of artificial intelligence that is trained on large datasets of text to understand, generate, and manipulate human language. In the

context of assistants, LLMs act as the brain that enables the assistant to understand questions, interact intelligently with users, and perform actions on your behalf. The Opsloom assistants integrate with some of the most well-known LLMs, including OpenAI's GPT models (which provide the basis for ChatGPT), Meta's LLaMa, and Anthropic's Claude.

Knowledge Base

A knowledge base is a repository of documents used to provide context for an assistant. During retrieval-augmented generation (RAG), these knowledge bases are treated as the authoritative source of truth. Documents in a knowledge base can be anything text-based: a training manual, technical documentation, or logs.

Retrieval-Augmented Generation

Retrieval-augmented generation (RAG) is a technique that provides assistants with an external knowledge base in order to improve the accuracy, relevance, and up-to-date nature of their responses. Behind the scenes, the process of setting up a RAG AI assistant involves converting documents into chunks of text, turning those chunks into vectors known as *embeddings*, and using vector similarity search to identify which chunk of information is best suited to answer a user's request.

Next page
Architecture